



BITS Pilani
Pilani Campus

Foundations of Conversational AI

Lecture 1 | Module 1

Course: AIMLCZG521

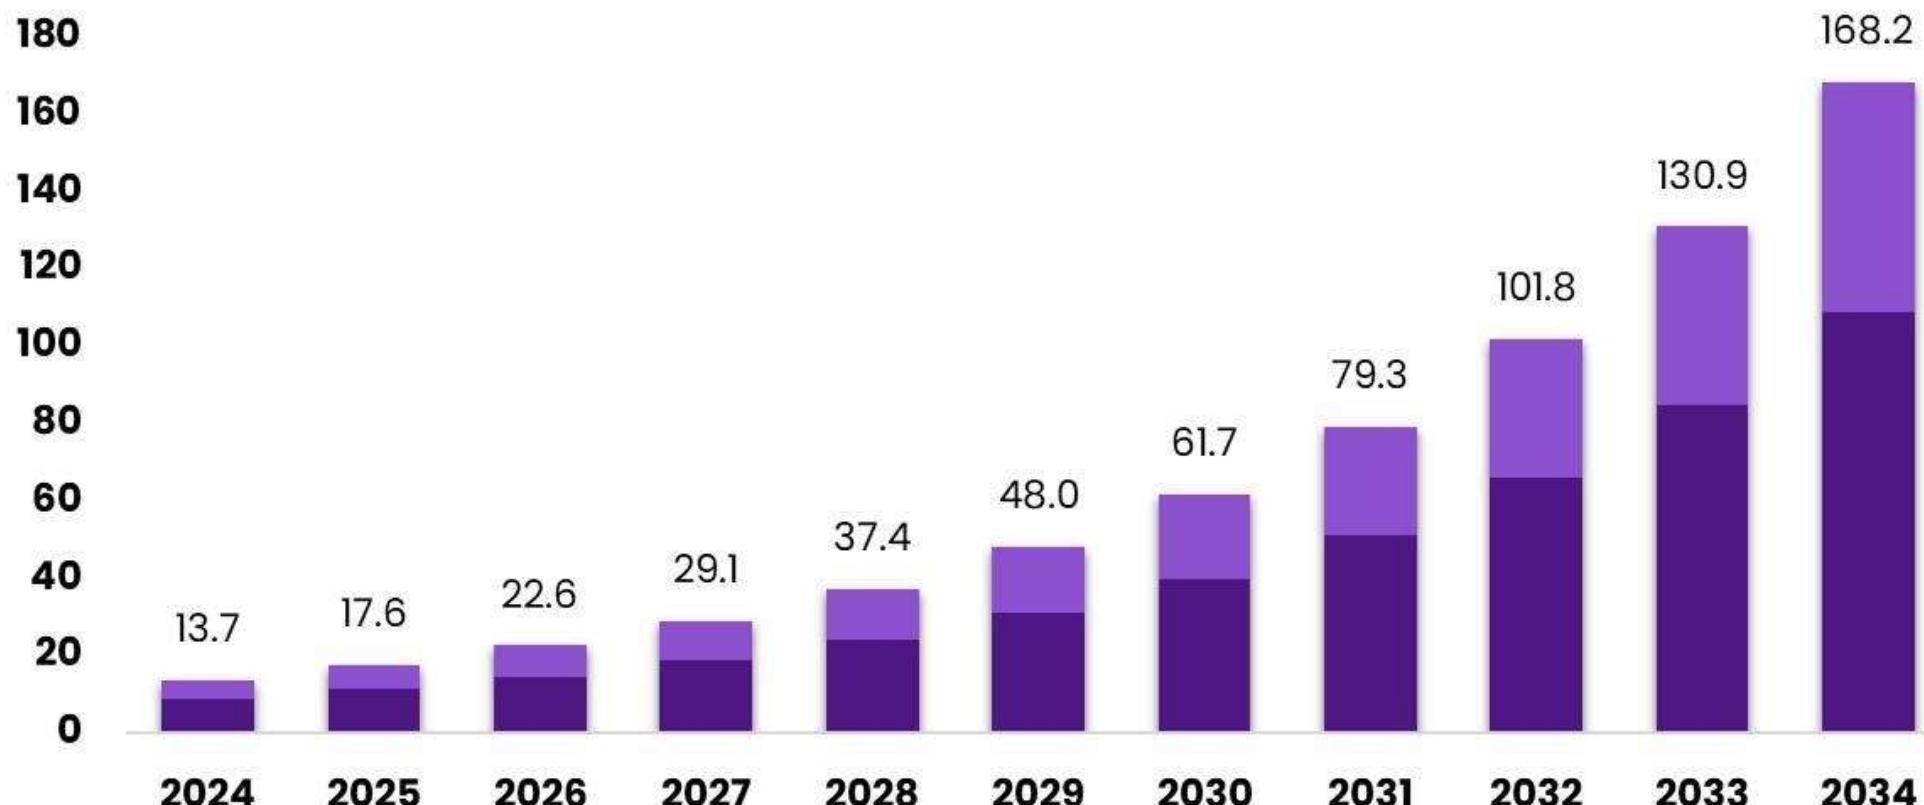
- Bhagath S

Conversational AI & Virtual Agents Market

Size, By Deployment, 2025-2034 (USD Billion)

■ On-Premises

■ Cloud Based



The Market will Grow
At the CAGR of:

28.5%

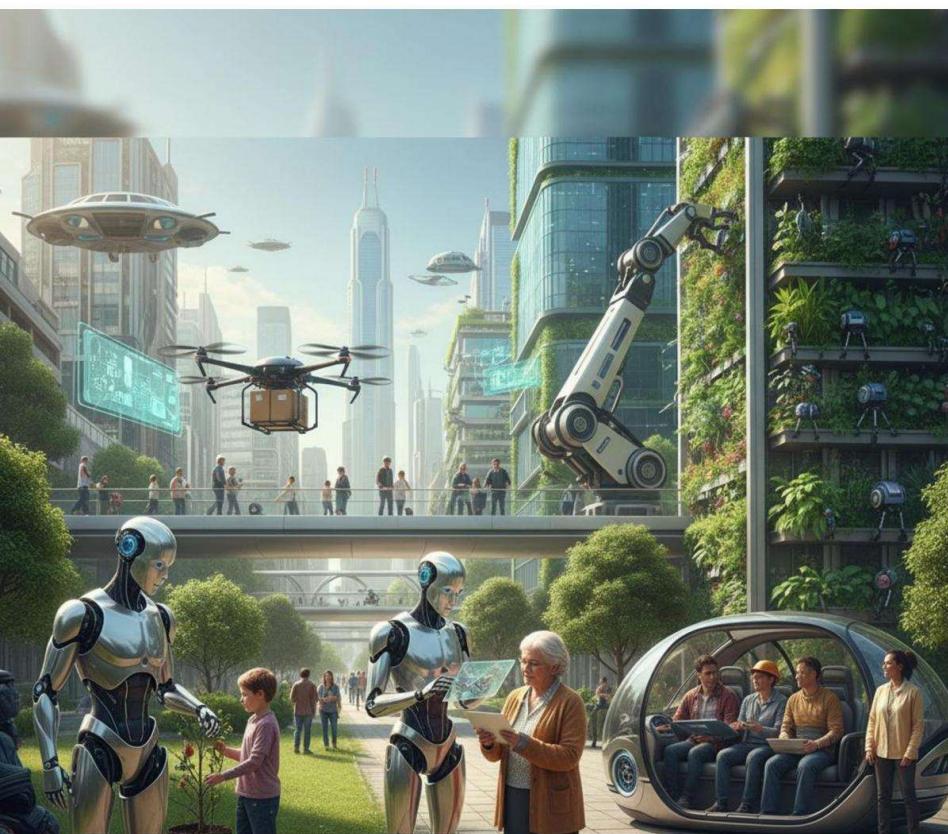
The Forecasted Market
Size for 2034 in USD:

\$168.2B

market.us
ONE STOP SHOP FOR THE REPORTS



What is Conversational AI?



Conversational AI refers to technologies that enable computers to understand, process, and respond to human language in natural, human-like ways through spoken or written interactions.

Core Capabilities

- ✓ • **Natural Language Understanding (NLU)**
Understanding user intent and entities
- ✓ • **Dialogue Management**
Managing conversation flow and context
- ✓ • **Natural Language Generation (NLG)**
Generating human-like responses

- ✓ • **Context Awareness**
Remembering conversation history
- ✓ • **Multi-turn Interaction**
Handling complex dialogues
- ✓ • **Personalization**
Adapting to user preferences

Examples in Daily Life

Voice Assistants: Alexa, Siri, Google Assistant

Customer Support: Banking chatbots, e-commerce help

Virtual Agents: Healthcare triage, HR assistants



The Evolution Journey (1960s - 2025)

Era	Technology	Capabilities	Limitations
1960s-1990s Rule-Based	Pattern matching, keyword detection (ELIZA, ALICE)	Simple Q&A, scripted responses	No learning, rigid, no context
2000-2010 Statistical ML	SVMs, CRFs, HMMs for NLU/NER	Intent classification, entity extraction	Limited context, hand-crafted features
2010-2017 Deep Learning	RNNs, LSTMs, Seq2Seq, Attention	End-to-end learning, better context	Data hungry, task-specific training
2017-2020 Transformers	BERT, GPT-2, T5 (pre-trained models)	Transfer learning, contextual embeddings	Still task-specific fine-tuning needed
2020-2023 LLMs & Gen AI	GPT-3, ChatGPT, Claude, PaLM	Few-shot learning, general purpose, fluent generation	Hallucinations, no real-time data, no actions
2023-2025 Agentic AI	LLMs + Tools + Memory + Planning	Execute actions, multi-step reasoning, tool use, autonomous workflows	Complex orchestration, scaling challenges, cost optimization needed

T-1 T-2 T-3. { Tool: "2" } ✓ { DOB: "22/MM/YYYY" }



Evolution Deep Dive: Early Stages

1. Rule-Based Systems (1960s-1990s)

Example: ELIZA (1966)

```
IF user_input contains "mother": RESPOND "Tell me more about your family" IF user_input contains  
"sad": RESPOND "Why do you feel sad?"
```

Limitations: No learning, brittle, couldn't handle variations, no real understanding



Evolution Deep Dive: Early Stages

2. Statistical ML Era (2000-2010)

Key Techniques:

- **Intent Classification:** SVM, Naive Bayes to classify user intent
- **Named Entity Recognition:** CRF, HMM for extracting entities
- **Dialogue State Tracking:** Markov models for conversation flow

Example Framework: Microsoft LUIS, IBM Watson (early versions)

Advance: Could learn from data, handle variations better

Limitations: Required feature engineering, limited context understanding



Evolution Deep Dive: Deep Learning Era

3. Deep Learning Revolution (2010-2017)

Key Architectures:

- **RNNs/LSTMs:** For sequence modeling, context retention
- **Seq2Seq Models:** Encoder-decoder for response generation
- **Attention Mechanism:** Focus on relevant parts of input

Example Frameworks: Rasa (open source), Google Dialogflow

Advance: End-to-end learning, better context, no hand-crafted features

Limitations: Data hungry, task-specific models, limited transfer learning



Evolution Deep Dive: Deep Learning Era

4. Transformer Era (2017-2020)

Breakthrough: "Attention is All You Need" (2017)

- **Pre-trained Models:** BERT, GPT-2, T5, RoBERTa
- **Transfer Learning:** Pre-train on massive data, fine-tune for tasks
- **Contextual Embeddings:** Word meaning depends on context

Impact on Conversational AI:

- Better intent classification with fewer examples
- Improved entity recognition
- More natural response generation



Evolution: Modern Era (2020-2025)

5. Large Language Models (2020-2023)

Game Changers: GPT-3 (2020), ChatGPT (2022), Claude, PaLM

- **Massive Scale:** Billions to trillions of parameters
- **Few-Shot Learning:** Learn tasks from examples in prompt
- **General Purpose:** One model for many tasks
- **Human-like Fluency:** Natural, contextual responses

Revolution for Conversational AI:

- No fine-tuning needed for many tasks
- Understand complex, multi-turn conversations
- Generate creative, contextual responses

But still limited: No real-time data, can't take actions, hallucinations



Evolution: Modern Era (2020-2025)

6. Agentic Conversational AI (2023-2025)

The Next Frontier: LLMs + Tools + Memory + Planning

- **Tool Use:** Can call APIs, search databases, execute code
- **Multi-step Reasoning:** Break down complex tasks
- **Memory Systems:** Remember user preferences, conversation history
- **Autonomous Actions:** Book appointments, send emails, create documents

This is what we'll build in this course!



Architecture Evolution: Then vs Now

Traditional Architecture (Pre-2020)

```
User Input ↓ Speech Recognition (if voice) ↓  
Natural Language Understanding • Intent  
Classification • Entity Extraction ↓ Dialogue  
Manager • State Tracking • Policy (rules/ML) ↓  
Natural Language Generation • Template-based •  
Rule-based ↓ Text-to-Speech (if voice) ↓  
Response
```

Frameworks: Rasa, Dialogflow, Microsoft Bot Framework, Amazon Lex

Modern Agentic Architecture (2023+)

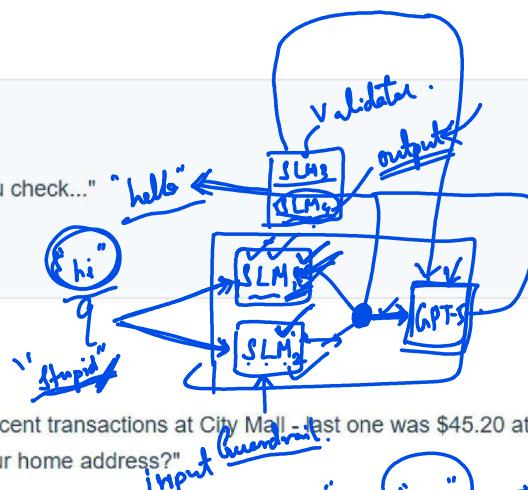
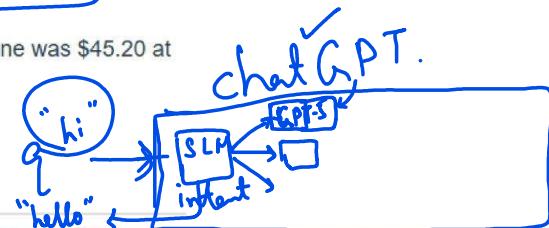
```
User Input ↓ LLM-based Understanding • Intent +  
Entities in one pass ↓ Agentic Orchestration  
Layer • Planning (break down task) • Tool  
Selection • Memory Retrieval ↓ Tool Invocation  
• API calls • Database queries • Code execution  
↓ Memory Update • Store context • Update user  
profile ↓ LLM-based Generation • Contextual  
response ↓ Safety & Validation ↓ Response
```

Frameworks: LangChain, LlamaIndex, AutoGPT, CrewAI

Use Case: Customer Support Evolution

Scenario: Banking Customer Support

LLM.
GPT-5

Approach	Capabilities	User Experience	Business Impact:
Rule-Based (2000s)	<ul style="list-style-type: none"> FAQ matching Keyword detection Fixed responses 	User: "I lost my card" Bot: "For lost card, press 1. For stolen, press 2" Rigid, frustrating	
Intent-Based (2015)	<ul style="list-style-type: none"> Intent classification Entity extraction Dialogue flow 	User: "I lost my card" Bot: "I'll help you block your card. Which card - Credit or Debit?" Better, but limited	
LLM-based (2023)	<ul style="list-style-type: none"> Natural conversation Context understanding Fluent responses 	User: "I think I left my card at the restaurant" Bot: "I understand. Let me help you block it temporarily while you check..." Natural, but no action	
Agentic AI (2025)	<ul style="list-style-type: none"> Everything above + Access banking system Execute transactions Multi-step actions 	User: "I lost my card at City Mall" Bot: "I've immediately blocked your card ending in 1234. I see recent transactions at City Mall - last one was \$45.20 at Starbucks 2 hours ago. Should I order a replacement card to your home address?" Proactive, action-oriented	



Components of Modern Conversational AI

1. Natural Language Understanding

Role: Understand what user wants

- Intent classification
- Entity extraction
- Sentiment analysis
- Context understanding

Modern approach: LLM-based, single pass

2. Dialogue Management

Role: Manage conversation flow

- State tracking
- Context maintenance
- Turn-taking
- Error handling

Modern approach: LLM reasoning + Memory systems



Components of Modern Conversational AI

3. Knowledge Access

Role: Retrieve relevant information

- Vector databases
- Semantic search
- RAG (Retrieval-Augmented Generation)
- Real-time data access

Modern approach: Embeddings + Hybrid retrieval

4. Action Execution

Role: Take actions on behalf of user

- Tool/function calling
- API integrations
- Database operations
- External service invocation

Modern approach: Agentic tool use



Components of Modern Conversational AI

5. Response Generation

Role: Generate natural responses

- Contextual generation
- Personality/tone
- Multi-modal output
- Structured responses

Modern approach: LLM generation with control

6. Memory Systems

Role: Remember user context

- Short-term (conversation)
- Long-term (user profile)
- Episodic memory
- Semantic memory

Modern approach: Vector + SQL hybrid



Framework Evolution for Conversational AI

Traditional Frameworks (2015-2020)

Framework	Approach	Use Case
Rasa	Intent + Entity + Dialogue Policies	Custom chatbots, on-premise deployment
Google Dialogflow	Intent-based, ML-powered NLU	Voice assistants, Google integrations
Microsoft Bot Framework	Intent + Entity + Dialogue Management	Enterprise bots, Azure integrations
Amazon Lex	Intent-based, Alexa-powered	Voice + text bots, AWS services

Key Shift: From intent-based dialogue systems to LLM-powered agentic systems with tool use and planning capabilities.



Framework Evolution for Conversational AI

Traditional Frameworks (2015-2020)

Key Shift: From intent-based dialogue systems to LLM-powered agentic systems with tool use and planning capabilities.

Modern Agentic Frameworks (2023-2025)



Framework Evolution for Conversational AI

Modern Agentic Frameworks (2023-2025)

Framework	Approach	Use Case
LangChain / LangGraph ✓	LLM orchestration + Tools + Memory + Workflows ✓	Complex agents, RAG, multi-step reasoning, production systems
LlamaIndex ✓	Data-centric, RAG-optimized pipelines	Document Q&A, knowledge bases, enterprise search
Semantic Kernel ✓	Microsoft's SDK for AI orchestration	Enterprise, .NET/Azure ecosystem integration
Haystack ✓	End-to-end NLP framework	Production RAG systems, search applications
AutoGen ✓	Multi-agent conversation framework	Complex workflows with agent collaboration



Course Content

The Paradigm Shift: Conversational AI has fundamentally changed from rule-based systems to intelligent, autonomous agents. This course prepares you for the **current and future** state of conversational AI.

Module 1: Foundations

- **Embeddings & Vector Search:** Modern NLU requires semantic understanding → vector embeddings enable finding relevant context and intents
- **Model Landscape:** Understanding LLMs (GPT-4, Claude, Gemini) that power modern conversational AI
- **Cost Engineering:** Real-world deployment requires optimizing token usage and model selection

Module 2: Core Building Blocks

- **Function Calling:** How conversational agents execute actions (book flights, query databases)
- **Memory Systems:** How agents remember conversation history and user preferences
- **RAG Pipelines:** How agents access knowledge bases to answer questions accurately



Course Content

Module 3: Autonomous Agents

- **Planning:** How agents break down complex user requests into steps
- **Multi-Agent Systems:** How specialized agents collaborate (customer support routing)
- **Evaluation:** Measuring quality of conversational AI systems

Module 4: Production Ecosystem

- **Security:** Protecting against prompt injection, PII leakage in conversations
- **Protocols (MCP, A2A):** Standard ways for agents to communicate
- **Ethics:** Ensuring fair, unbiased conversational AI systems



NLP Foundations for Conversational AI

Core NLP concepts that directly impact conversational agent design and performance

1. Tokenization - The Foundation

What: Breaking text into subword units (tokens) that LLMs process

Why it matters for Conversational AI:

- **Cost:** API pricing is per token → impacts conversation economics
- **Context Window:** Limits conversation length (e.g., 200K tokens = ~150K words)
- **Latency:** More tokens = slower response in conversations
- **Quality:** Tokenization affects understanding of domain-specific terms





NLP Foundations for Conversational AI

2. Context Windows - Conversation Memory

What: Maximum text length (in tokens) an LLM can process at once

Current State (2025):

- GPT-4 Turbo: 128K tokens (~96K words) - standard production use
- Claude 3.5 Sonnet: 200K tokens (~150K words) - extended conversations
- Gemini 1.5 Pro: 1M tokens (~750K words) - entire codebases
- Emerging models: 2M+ tokens for specialized use cases

Challenge: "Lost in the middle" - models struggle with info in long contexts

Solution: RAG + Memory systems (covered in Module 2)



Tokenization: How It Works

Common Tokenization Algorithms

1. Byte-Pair Encoding (BPE) - Used by GPT models
2. SentencePiece - Used by many multilingual models
3. WordPiece - Used by BERT and similar models

Example: How Text Becomes Tokens

Text	Tokens	Count
"Hello World"	["Hello", " World"]	2 tokens
"artificial intelligence"	["art", "ificial", " intelligence"]	3 tokens
"GPT-4"	["G", "PT", "-", "4"]	4 tokens
"Book a flight to NYC"	["Book", " a", " flight", " to", " NYC"]	5 tokens



Tokenization: How It Works

Impact on Conversational AI

Scenario: Customer support conversation (2025 typical costs)

- Average conversation: 40-60 turns (user + agent exchanges)
- Average tokens per turn: 15-25 tokens
- Total per conversation: ~800-1,200 tokens
- Cost with GPT-4o: ~\$0.01-0.03 per conversation
- Cost with GPT-3.5 Turbo: ~\$0.002-0.005 per conversation
- At 10,000 conversations/day with GPT-4o: \$100-300/day

Key insight: Model selection and optimization can reduce costs by 10-20x. We'll cover these techniques in Lecture 11.



Hands-On: Byte-Pair Encoding

Understanding Tokenization with Code

We'll explore how text is converted to tokens using the tiktoken library

Demo Components:

- **Text to Tokens:** See how different texts are tokenized
- **Token Counting:** Calculate tokens for a sample conversation
- **Cost Analysis:** Understand the economic implications
- **Model Comparison:** Tokenization differences across models



Transformers & LLMs: The Brain of Modern Conversational AI

You learned about **transformers** in your Deep Learning course. Let's connect that to **Conversational AI**.

What Transformers Gave Us

Technical Capabilities

- **Self-Attention**: Understand relationships between words
- **Contextual Understanding**: Word meaning depends on context
- **Long-Range Dependencies**: Connect ideas across long text
- **Parallel Processing**: Faster than RNNs

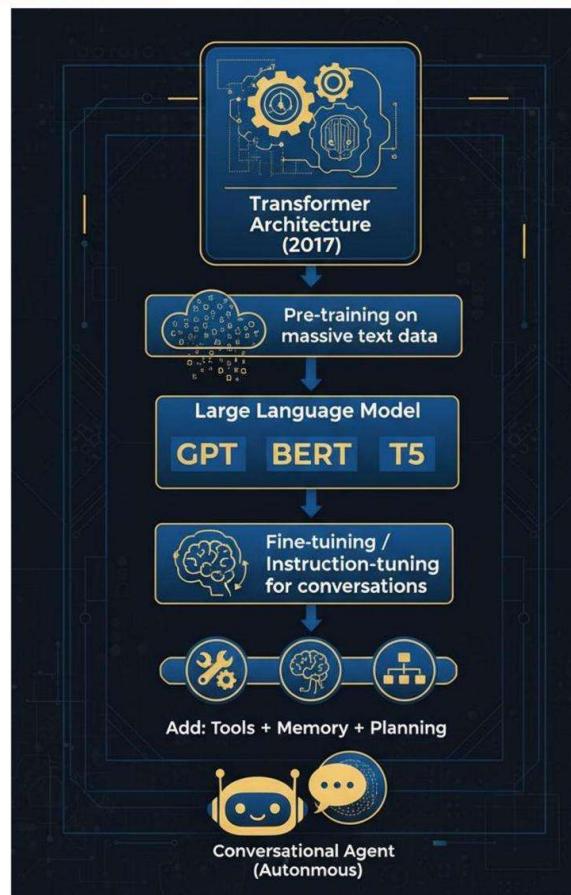
Impact on Conversations

- **Multi-turn coherence**: Remember earlier in conversation
- **Intent understanding**: Grasp user goals from context
- **Natural responses**: Generate human-like replies
- **Ambiguity handling**: Resolve unclear requests



Transformers & LLMs: The Brain of Modern Conversational AI

From Transformer → LLM → Conversational Agent



Key Insight: LLM is the "brain" but needs additional systems (tools, memory, planning) to become a fully functional conversational agent.



LLM Capabilities & Limitations for Conversational AI

✓ What LLMs Do Well

- **Natural conversation:** Human-like dialogue
- **Intent understanding:** Grasp user goals
- **Context retention:** Multi-turn conversations
- **Response generation:** Fluent, contextual replies
- **Summarization:** Condense information
- **Entity extraction:** Identify key information
- **Sentiment analysis:** Understand user emotions

✗ LLM Limitations

- **No real-time data:** Training cutoff date
- **Hallucinations:** Generate plausible but false info
- **Can't take actions:** No access to external systems
- **No memory:** Forget after context window
- **Calculation errors:** Struggle with math
- **Consistency:** Different answers to same question
- **No verification:** Can't check facts



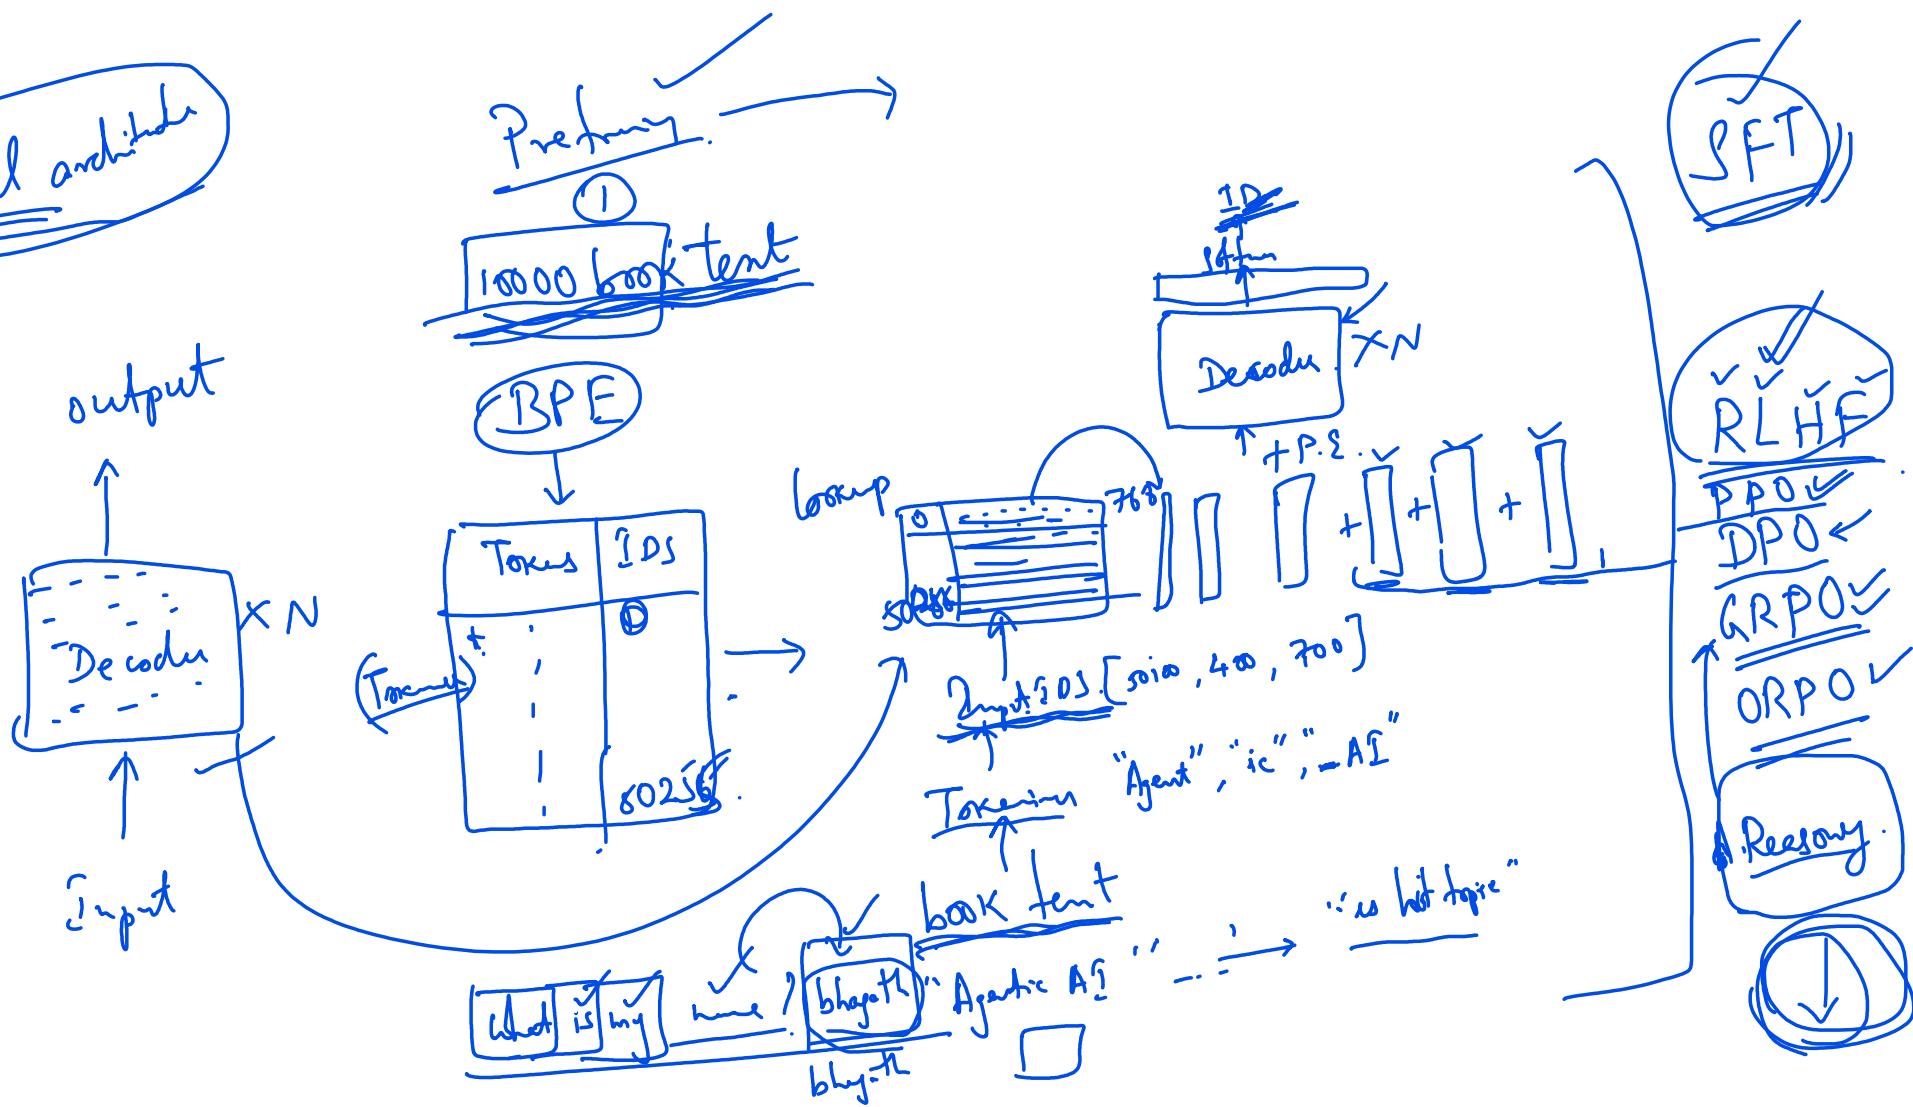
LLM Capabilities & Limitations for Conversational AI

Why We Need Agentic Architecture

Solution: Augment LLMs with external capabilities

LLM Limitation	Agentic Solution	Example
No real-time data	Tool calling (APIs, search)	Check current flight prices
Hallucinations	RAG (Retrieval-Augmented Generation)	Ground answers in company docs
Can't take actions	Function calling	Book appointment, send email
No memory	External memory systems	Remember user preferences
Calculation errors	Calculator tool, code execution	Precise math computations

Model architecture





Agentic Conversational AI: Request → Response Pipeline

How a modern conversational agent processes user input

Step	Details
1. REQUEST	<ul style="list-style-type: none">User input received and validatedParse user intentExtract entitiesInput sanitization (security)
2. ROUTING	<ul style="list-style-type: none">Intent classification & task decompositionClassify intent typeDetermine required toolsRoute to appropriate sub-agent
3. REASONING	<ul style="list-style-type: none">Planning & decision making with LLMBreak down into stepsIdentify information gapsPlan tool call sequence
4. TOOL INVOCATION	<ul style="list-style-type: none">External API calls & actionsExecute API callsDatabase queriesCode execution
5. MEMORY	<ul style="list-style-type: none">Context storage & retrievalStore conversation historyUpdate user preferencesMaintain session state
6. SAFETY	<ul style="list-style-type: none">Guardrails & output validationCheck for toxic contentVerify factual accuracyPII redaction
7. RESPONSE	<ul style="list-style-type: none">Final output delivered to userNatural language generationFormat for channel (text/voice/UI)Send to user



Agentic Conversational AI: Request → Response Pipeline

How a modern conversational agent processes user input

Step	Details
1. REQUEST	<ul style="list-style-type: none">User input received and validatedParse user intentExtract entitiesInput sanitization (security)
2. ROUTING	<ul style="list-style-type: none">Intent classification & task decompositionClassify intent typeDetermine required toolsRoute to appropriate sub-agent
3. REASONING	<ul style="list-style-type: none">Planning & decision making with LLMBreak down into stepsIdentify information gapsPlan tool call sequence
4. TOOL INVOCATION	<ul style="list-style-type: none">External API calls & actionsExecute API callsDatabase queriesCode execution

Step	Details
5. MEMORY	<ul style="list-style-type: none">Context storage & retrievalStore conversation historyUpdate user preferencesMaintain session state
6. SAFETY	<ul style="list-style-type: none">Guardrails & output validationCheck for toxic contentVerify factual accuracyPII redaction
7. RESPONSE	<ul style="list-style-type: none">Final output delivered to userNatural language generationFormat for channel (text/voice/UI)Send to user



Lifecycle Example: Banking Conversational Agent

User Request: "What's my account balance and recent transactions?"

Stage	What Happens	Stage	What Happens
1. Request	<ul style="list-style-type: none">Intent: account_inquiryEntities: {type: ["balance", "transactions"]}Validate: User authenticated	4. Tool Invocation	<ul style="list-style-type: none">Call: get_balance(user_id) → Returns: \$5,432.10Call: get_transactions(user_id, limit=5)Returns: [Transaction1, Transaction2, ...]
2. Routing	<ul style="list-style-type: none">Route to: Banking AgentRequired tools: [get_balance, get_transactions]Permission check: User has access	5. Memory	<ul style="list-style-type: none">Store: User asked about balance at 2:30 PMUpdate: User preference for transaction detailsContext: Maintain for follow-up questions
3. Reasoning	<ul style="list-style-type: none">Plan:<ol style="list-style-type: none">Fetch account balanceFetch recent transactions (last 5)Format response with both	6. Safety	<ul style="list-style-type: none">Validate: Balance and transactions are for correct userCheck: No PII exposure in logsVerify: Response contains only authorized info



Lifecycle Example: Banking Conversational Agent

User Request: "What's my account balance and recent transactions?"

Stage	What Happens
7. Response	"Your current balance is \$5,432.10. Here are your recent transactions: 1. Starbucks - \$4.50 (Today) 2. Amazon - \$89.99 (Yesterday) 3. Shell Gas - \$45.00 (Yesterday) 4. Grocery Store - \$123.45 (2 days ago) 5. Netflix - \$15.99 (3 days ago) Would you like more details on any of these?"



Exercise: Map the Agent Lifecycle

Scenario:

"Find me a good Italian restaurant near my office that's open tonight and make a reservation for 2 at 7 PM"

Map This Request to the 7-Stage Lifecycle



Exercise: Map the Agent Lifecycle

Scenario:

"Find me a good Italian restaurant near my office that's open tonight and make a reservation for 2 at 7 PM"

Stage	What Needs to Happen?
1. Request	What entities need to be extracted from user input?
2. Routing	Which tools or agents are needed for this task?
3. Reasoning	What's the step-by-step plan to accomplish this?
4. Tool Invocation	What specific API calls need to be made?
5. Memory	What information should be stored for future use?
6. Safety	What validations are required before taking action?
7. Response	How should the agent communicate with the user?



Building a Simple Conversational Agent

Demonstration: Weather Information Agent

A practical example using native OpenAI API in Python

Implementation Stages

- Baseline:** LLM without tools - sees limitations
- Tool Definition:** Define weather function schema
- Tool Selection:** LLM decides when to use tool
- Execution:** Call actual weather API
- Response Generation:** LLM creates natural answer

Example Interaction

User: "What's the weather like in Mumbai?"

Agent:

- Extracts location: "Mumbai"
- Calls weather API
- Receives: 32°C, Humid, Partly Cloudy
- Responds: "It's currently 32°C in Mumbai with partly cloudy skies and high humidity."



Protocol Landscape for Conversational Agents

As conversational agents become more prevalent, they need **standardized ways to communicate** with tools, data, each other, and UIs.

Why Protocols Matter

Without Standards

- Every vendor has custom API
- Agents can't interoperate
- Vendor lock-in
- Duplicated integration effort

With Standards

- Plug-and-play integrations
- Multi-agent collaboration
- Vendor flexibility
- Ecosystem growth



Protocol Landscape for Conversational Agents

The Protocol Landscape (2025)

Protocol	Purpose	Use Case
MCP (Model Context Protocol)	Standardize how LLMs connect to data sources and tools	Agent accessing databases, files, APIs consistently
OpenAI Assistant API	Built-in tools, file access, code interpreter	Rapid agent development with OpenAI infrastructure
LangGraph Protocol	Graph-based agent workflows and state management	Complex multi-step reasoning, agent coordination
Custom REST/GraphQL APIs	Traditional service integration	Enterprise systems, legacy integrations

Note: The protocol landscape is rapidly evolving. Standards like MCP are emerging, while many production systems still use custom APIs. We'll explore these in detail in Lectures 13-14.



Production Concerns for Conversational AI

Building conversational agents that work in development is one thing. Building them to work **reliably at scale** in production is another.

1. Observability

What to track:

- Conversation flows
- Tool invocations
- Token usage per conversation
- Response latencies
- Error traces

Tools: LangSmith, Arize Phoenix, OpenTelemetry

2. Cost Management

Optimization strategies:

- Prompt caching (50-90% cost reduction)
- Model routing (use smaller models when appropriate)
- Token budgets per user/session
- Efficient retrieval (reduce context size)

Example: A customer support conversation can cost \$0.02-0.10 depending on optimization



Production Concerns for Conversational AI

3. Latency Budgets

Target latencies:

- Chat responses: < 2 seconds
- Tool execution: < 5 seconds
- Complex reasoning: < 30 seconds

Key: Set user expectations with progress indicators

4. Safety & Security

Defense layers:

- Input validation (prompt injection defense)
- PII detection and redaction
- Output filtering (toxicity, hallucinations)
- Human-in-the-loop for critical actions

Principle: Never rely on a single safety layer

These must be considered from day one—we'll cover each in detail in Lectures 11-12



Key Takeaways

1. Evolution of Conversational AI

From rule-based (1960s) → ML (2000s) → Deep Learning (2010s) → Transformers (2017) → LLMs (2020) → Agentic AI (2023+)

2. Architecture Transformation

Traditional: Intent → Dialogue Manager → Template Response

Modern: LLM + Tools + Memory + Planning + Safety

4. LLMs are the Brain

But they need external systems for tools, memory, real-time data, and action execution

5. Agent Lifecycle

Request → Routing → Reasoning → Tool Invocation → Memory → Safety → Response

3. Components Matter

Modern systems need: NLU, Dialogue Management, Knowledge Access, Action Execution, Response Generation, Memory

6. Production is Different

Observability, cost optimization, latency management, and safety are non-negotiable



Key Takeaways

The Foundation is Set!

You now understand what Conversational AI is, how it evolved, and the architecture of modern agentic systems.

Next: We'll cover these in detail, starting with embeddings and vector search.



Coming Up Next

Lecture 2: Embeddings, Vector Search & Hybrid Retrieval

What We'll Cover

- Deep dive into embeddings and semantic search
- Vector database architecture (HNSW, ANN)
- BM25 + Dense retrieval + RRF (hybrid search)
- Building a practical retrieval system
- Evaluation metrics for retrieval

How to Prepare

- **Pre-read:** "Dense Passage Retrieval for Open-Domain QA" (Karpukhin et al., 2020)
- **Set up:** Python environment with OpenAI/Anthropic API keys
- **Install:** Libraries: openai, chromadb, rank-bm25
- **Review:** Today's slides on Canvas



Questions & Discussion

Topics for Discussion:

- Conversational AI use cases in your organization
- Evolution and current state of the field
- Technical challenges you foresee
- Clarifications on any concepts
- Course logistics and expectations

Thank you!

See you in Lecture 2! 